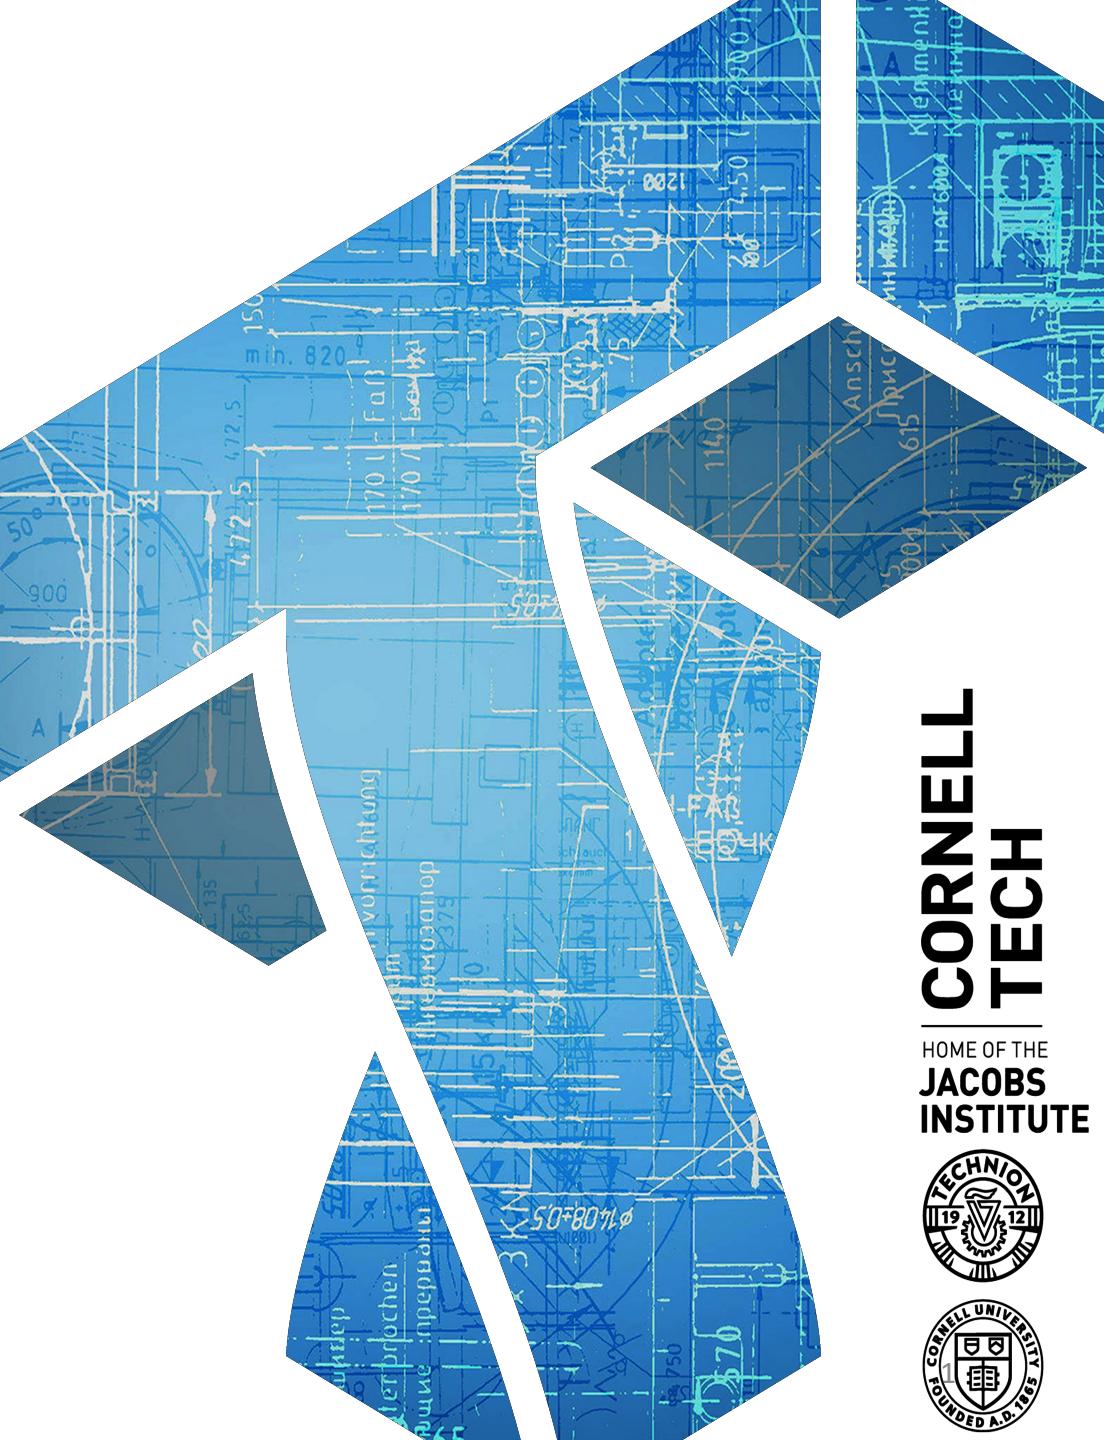


# CS 5830

# Cryptography



**CORNELL  
TECH**

HOME OF THE  
**JACOBS**  
**INSTITUTE**

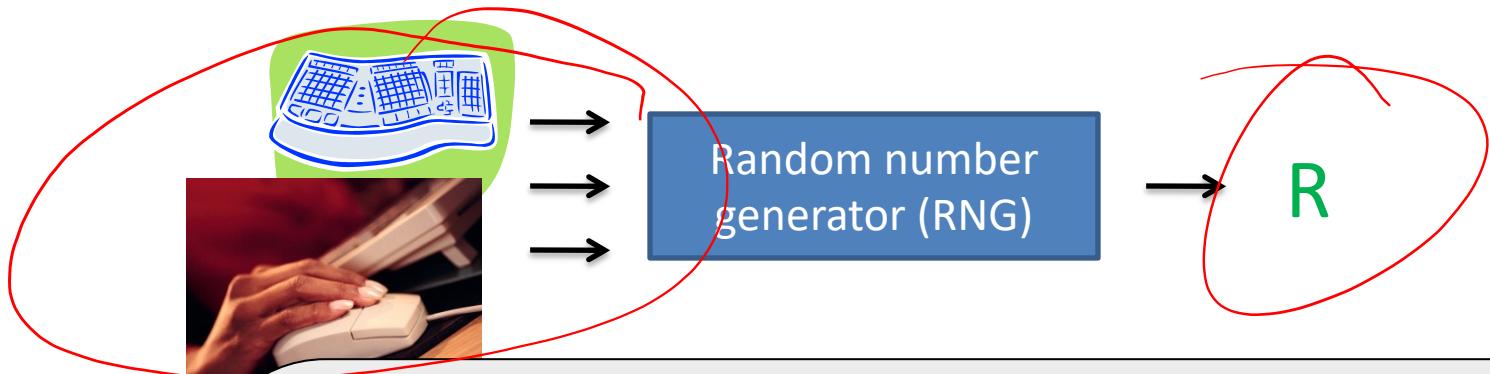


# Last time: passwords

- Passwords widely used in practice
- Never store passwords in clear, should at least be salted and slow hashed (PBKDF)
  - Memory-hard hashes are state-of-the art (scrypt)
- Password breach alerting increasingly used to warn users about exposed passwords
  - Cloudflare product (that I helped with): <https://blog.cloudflare.com/privacy-preserving-compromised-credential-checking/>
- Password-based encryption should derive key for AEAD scheme using PBKDF
  - Beware of what you see on internet about applied crypto

# Today: random number generation

- Cryptography relies on random numbers. What have we been using RNGs for?
  - Secret key generation
  - IVs for AEAD modes (CTR mode, CBC mode, AES-GCM nonces)
  - Salts for password hashing



## Random

[Wagner, G.

[Guttermar

[Guttermar

[Dorrendor

[Woolley et

[Bello 2008

[Mueller 20

[Abeni et al.

[Yilek et al. 2009]

[Heninger et al. 2012]

[Everspaugh et al. 2015]

```
MD_Update(&m,buf,j);
...
MD_Update(&m,buf,j); /* purify complains */
```

These lines of code commented out from OpenSSL random number generator code (`md_rand.c`) to **address complaints by security tools**  
**Purify and Valgrind**

Only the process ID (PID) was used as input to RNG.

It took a ~2 years for the bug to be (publicly) discovered!

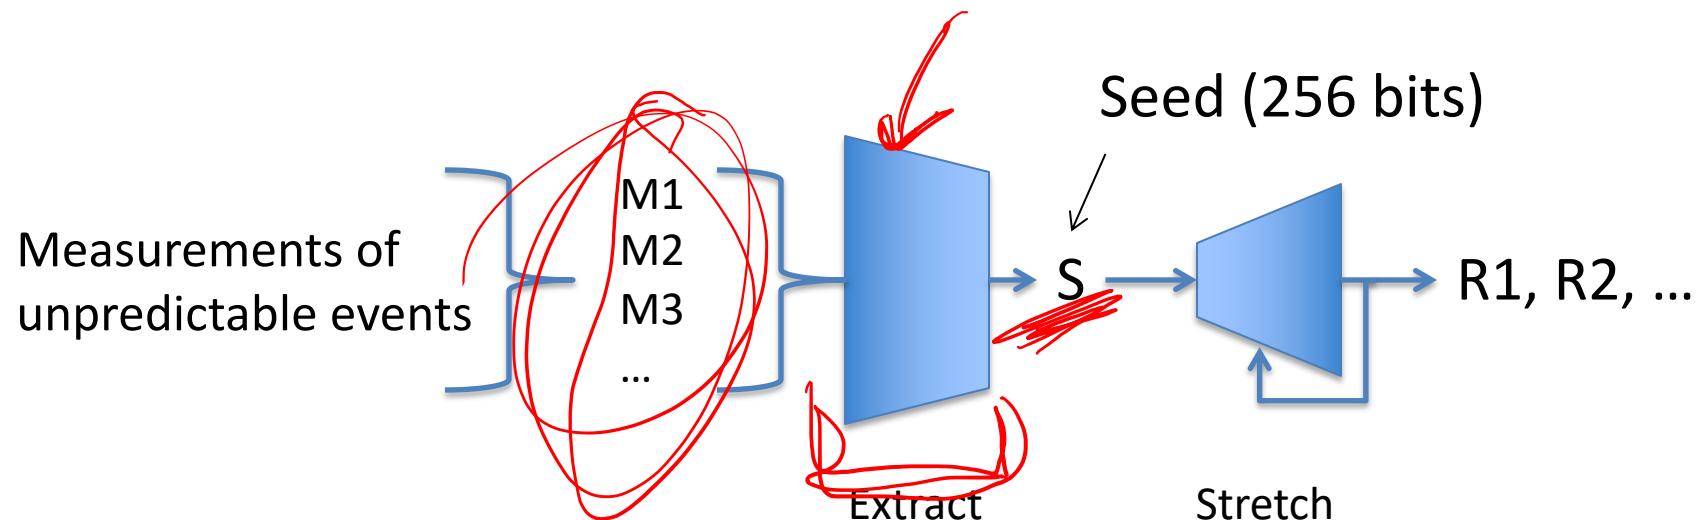
Debian OpenSSL bug lead to small set of possible **R**

# Cryptographically strong randomness

- Must be maximally unpredictable from adversary's perspective
- This means (computationally) indistinguishable from uniform bit string of same length
- “True” randomness vs. cryptographic randomness
  - Typically false dichotomy in practice

# RNG pipelines

1. Entropy gathering
2. Extracting from measurements a cryptographically strong value called seed
3. Using seed to deterministically produce pseudorandom values



# Extract-and-Expand approaches: HKDF

Hash-based key derivation function (HKDF)

- Derive more key material from one key
- Extract key from non-uniform key material

Uses HMAC with underlying hash function (SHA256)

HKDF(L,K,salt,info):

prk <- HMAC(salt,K)

$K_0$  <- empty byte string

m <- ceil(L/hashLen)

For i = 1 to m

$K_i$  <- HMAC(prk, $K_{i-1}$  || info ||  $<i+1>$ )

Return truncate<sub>L</sub>( $K_1$  || ... ||  $K_m$ )

# Entropy sources



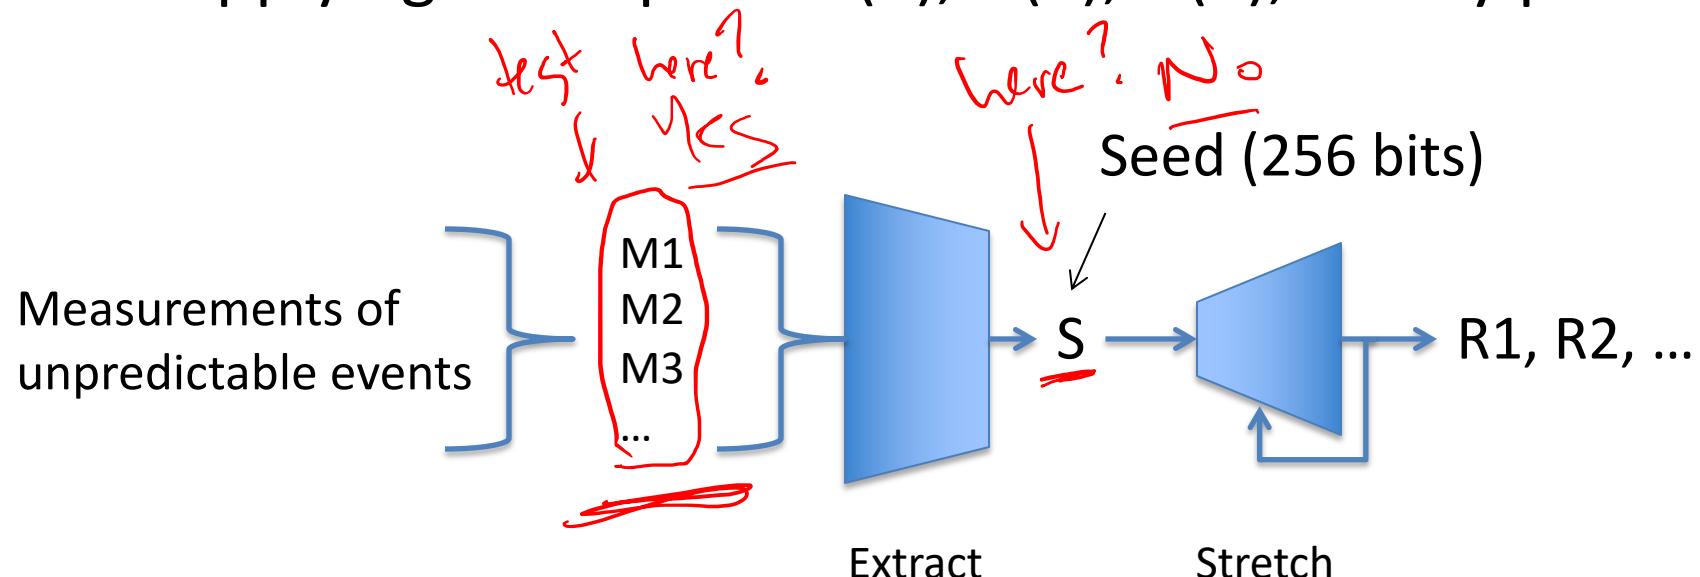
- Timing and description of various events
  - keyboard presses and timing
  - file/network interrupts
  - mouse movements
- Hardware RNGs
  - Intel RNG has custom hardware for generating unpredictable bits using thermal noise
  - RDRAND instruction
- Health tests

*Richard tests*

# How can we test security of an RNG?

- Statistical tests
  - Diehard tests, NIST 800-22 tests
- Apply to inputs or to outputs?
  - Careful of applying to outputs:  $H(0), H(1), H(2), \dots$  may pass tests

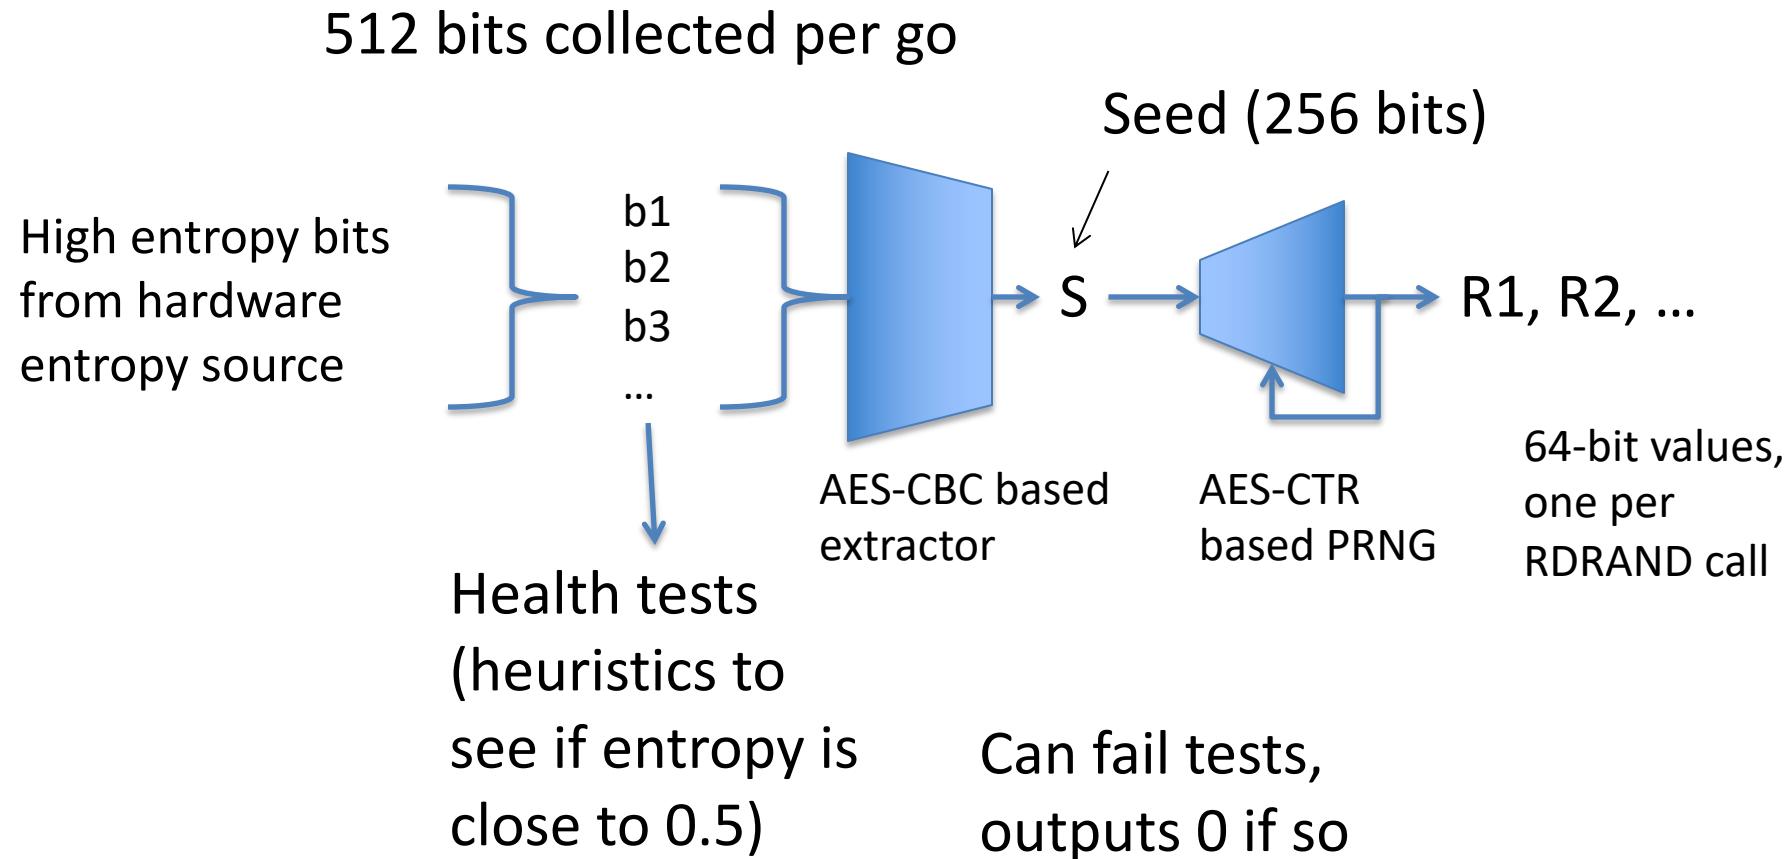
$$b_1, b_2, \dots, b_{256} \leftarrow H(0) \\ \sum_i b_i - \sum(b_i) \leq S$$



# Testing randomness

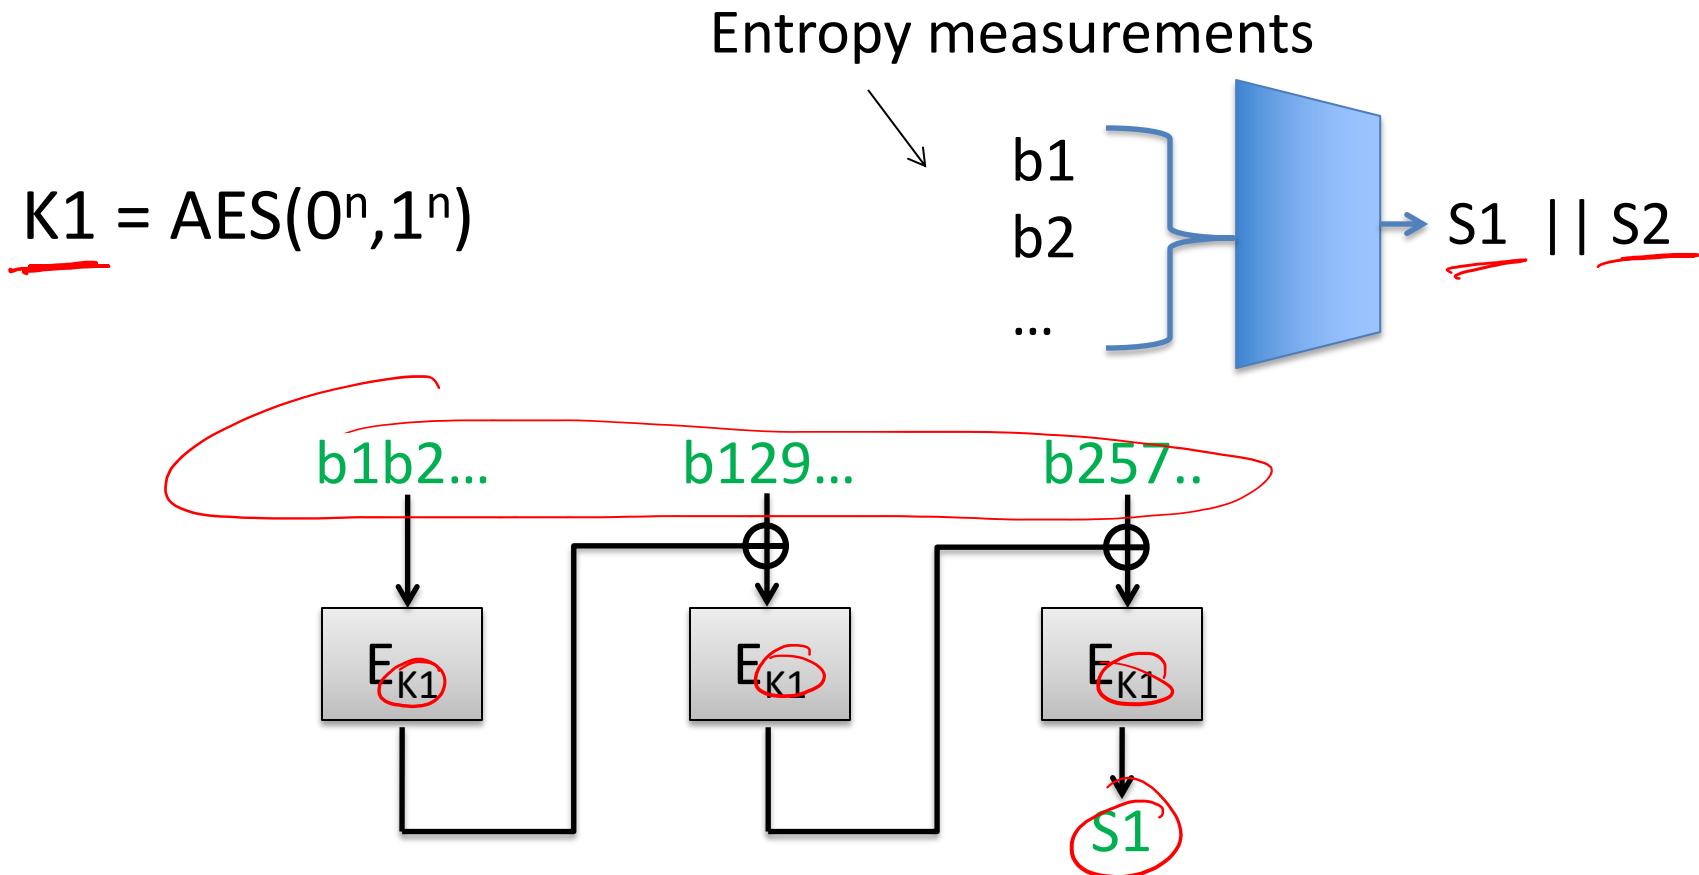
- “However, no set of statistical tests can absolutely certify a generator as appropriate for usage in a particular application, i.e., statistical testing cannot serve as a substitute for cryptanalysis.” from NIST 800-22  
(<https://nvlpubs.nist.gov/nistpubs/legacy/sp/nistspecialpublication800-22r1a.pdf>)
- Tests can help catch problems, but best bet is to try to build attacks against specific RNGs

# Intel RNG system



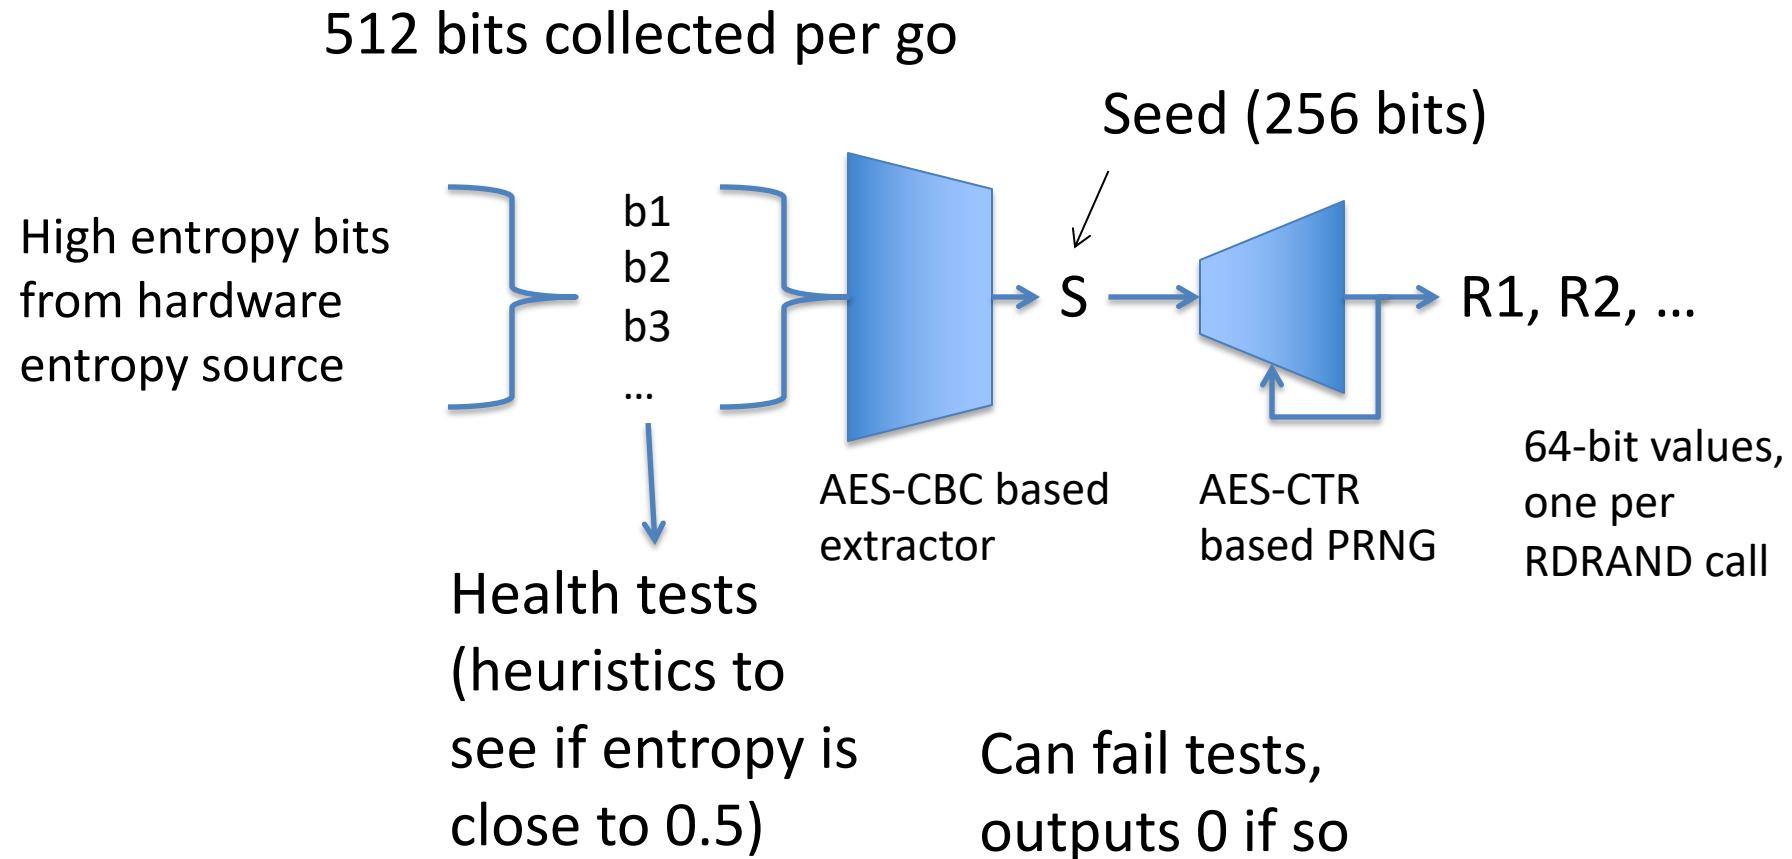
Good writeup: <http://eprint.iacr.org/2014/504.pdf>

# AES CBC MAC as an extractor



Repeat process of collecting entropy  
values and CBC-MACing to get  $S_2$

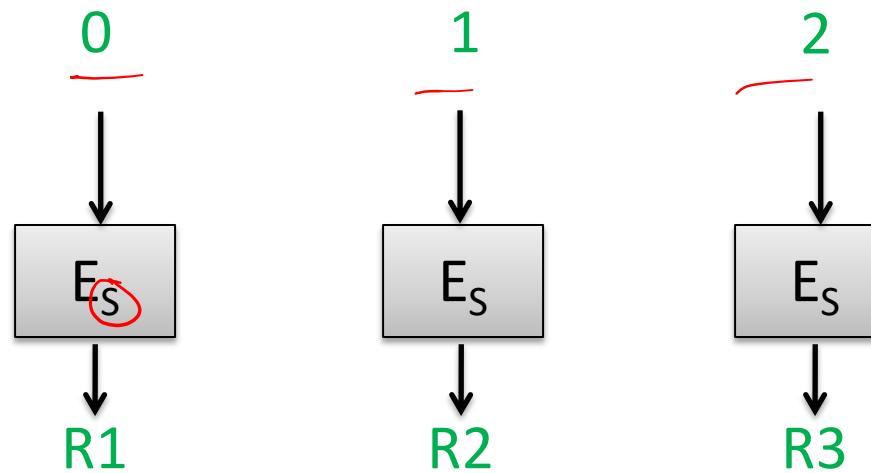
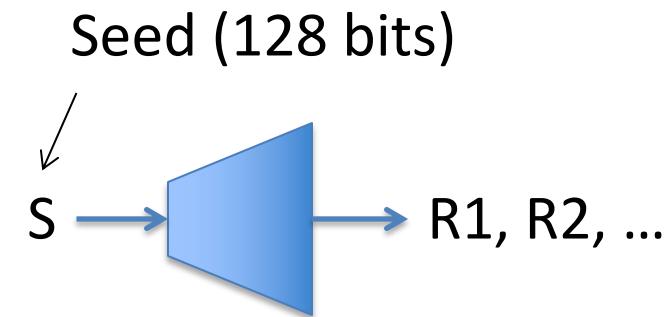
# Intel RNG system



Good writeup: <http://eprint.iacr.org/2014/504.pdf>

# AES CTR mode as PRG

AES-CTR( $S$ )  $\rightarrow R_1, R_2, R_3 \dots$

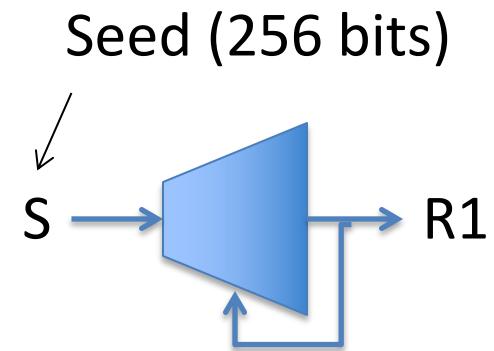


# AES CTR mode in Intel RNG

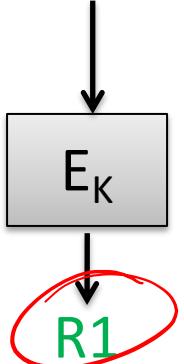
$\text{AES-CTR}(K, IV, S) \rightarrow R1, K', IV'$

$K, IV$  initially all zeros

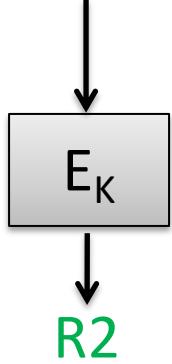
$S = S1 \parallel S2$  (128 bits each)



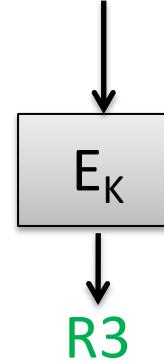
IV



IV + 1



IV + 2



R1 output to caller  
of instruction

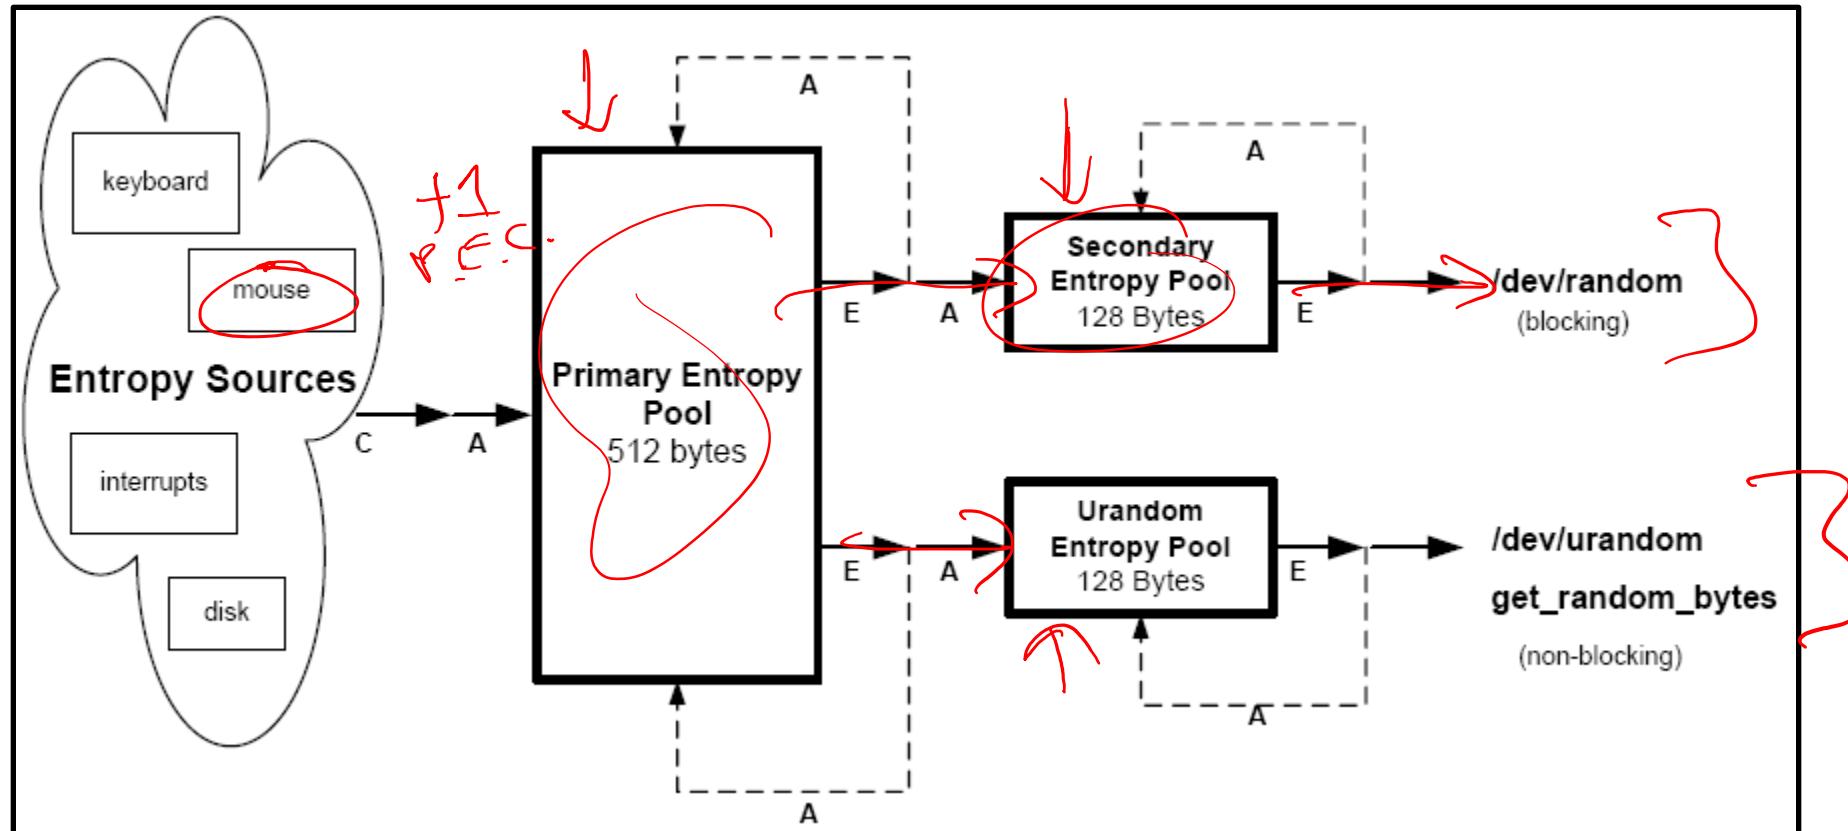
$$\oplus \quad S1 \\ K'$$

$$\oplus \quad S2 \\ IV'$$

# Linux /dev/(u)random

Linux random number generator (2500 lines of undocumented code)

Diagram from [Guterman, Pinkas, Reinman 2006]



Primary entropy pool feeds into other entropy pools only when 192 bits of "entropy" are estimated. Favors /dev/random

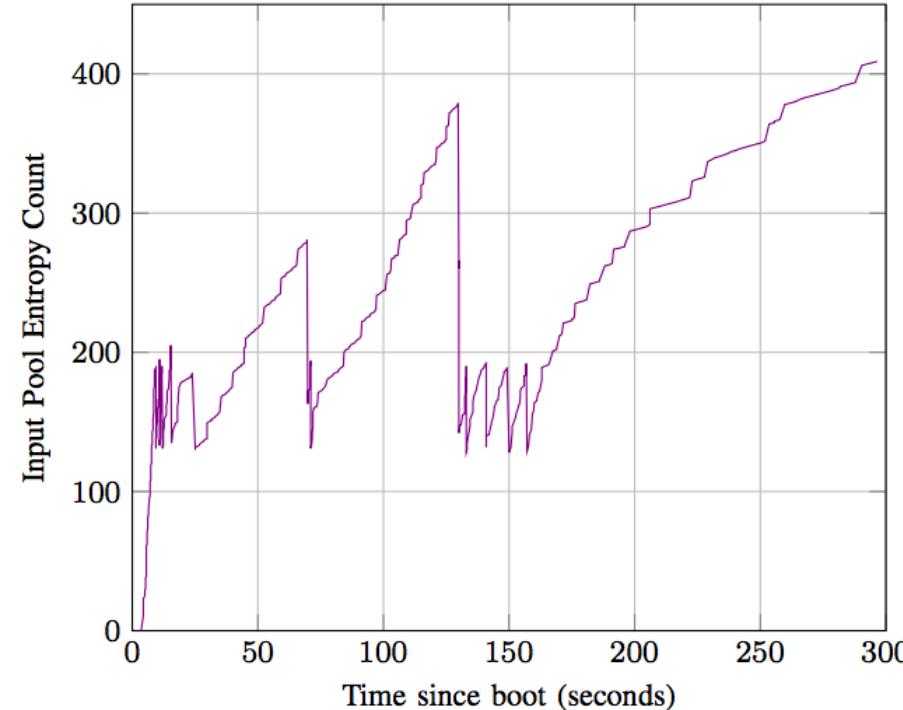
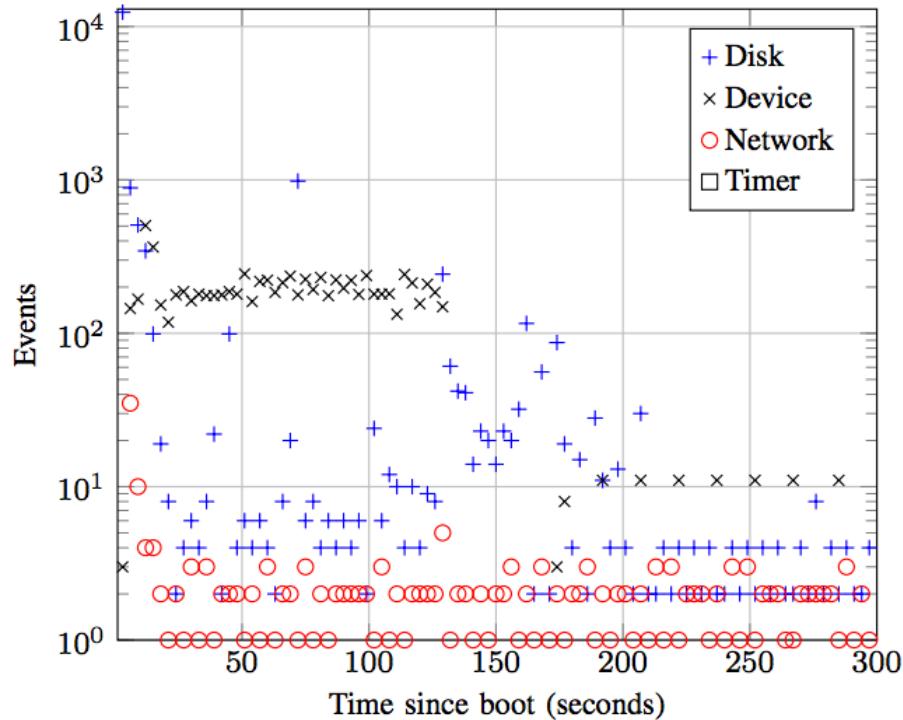
# Questions

- Does /dev/(u)random collect sufficient entropy during boot?
  - Stamos, Becherer, and Wilcox conjecture not in virtualized environments (BlackHat 2009)
- What happens when a full-state snapshot is resumed?

We carefully instrumented Linux kernel to track entropy accumulation

<http://pages.cs.wisc.edu/~ace/papers/not-so-random.pdf>

# Entropy accumulation during boot of Linux VM within VMWare



Our analysis suggests that, after first use of /dev/urandom during boot, entropy is sufficient to prevent attacks

# Boot-time entropy holes

First read from /dev/urandom before any entropy inputs.

Output is always: 0x22DAE2A8 862AAA4E

Combined with cycle counter to seed stack canary on init process

Not clear how to exploit directly

Embedded systems also exhibit boot-time entropy holes:

urandom entropy pool not updated long into boot

ssh keys generated on first boot --- broken!

<https://factorable.net/weakkeys12.extended.pdf>

# Questions

- Does /dev/(u)random collect sufficient entropy during boot?
  - Stamos, Becherer, and Wilcox conjecture not in virtualized environments (BlackHat 2009)
- What happens when a full-state snapshot is resumed?

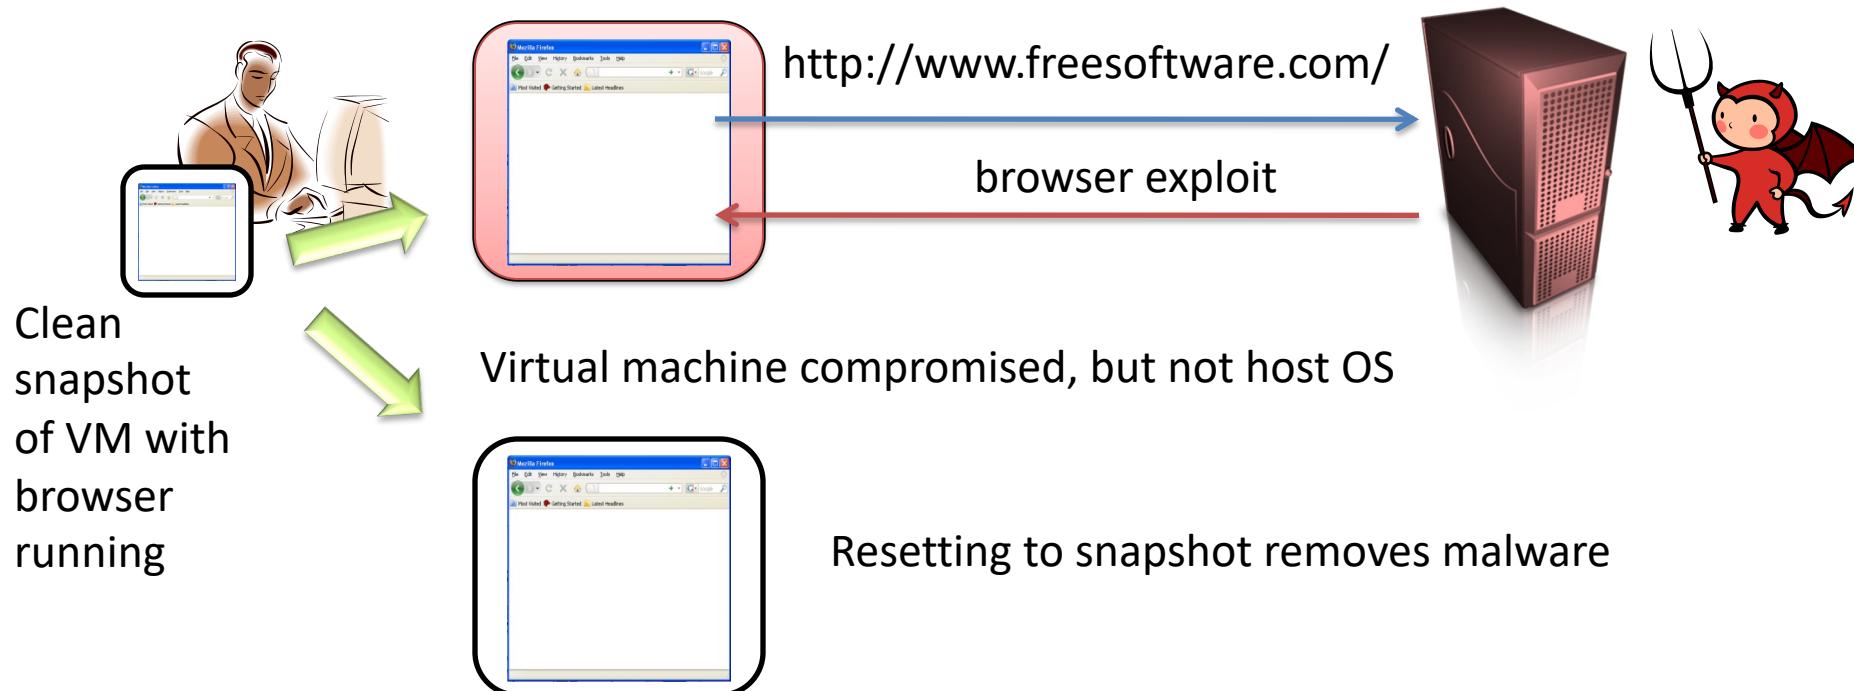
We carefully instrumented Linux kernel to track entropy accumulation

<http://pages.cs.wisc.edu/~ace/papers/not-so-random.pdf>

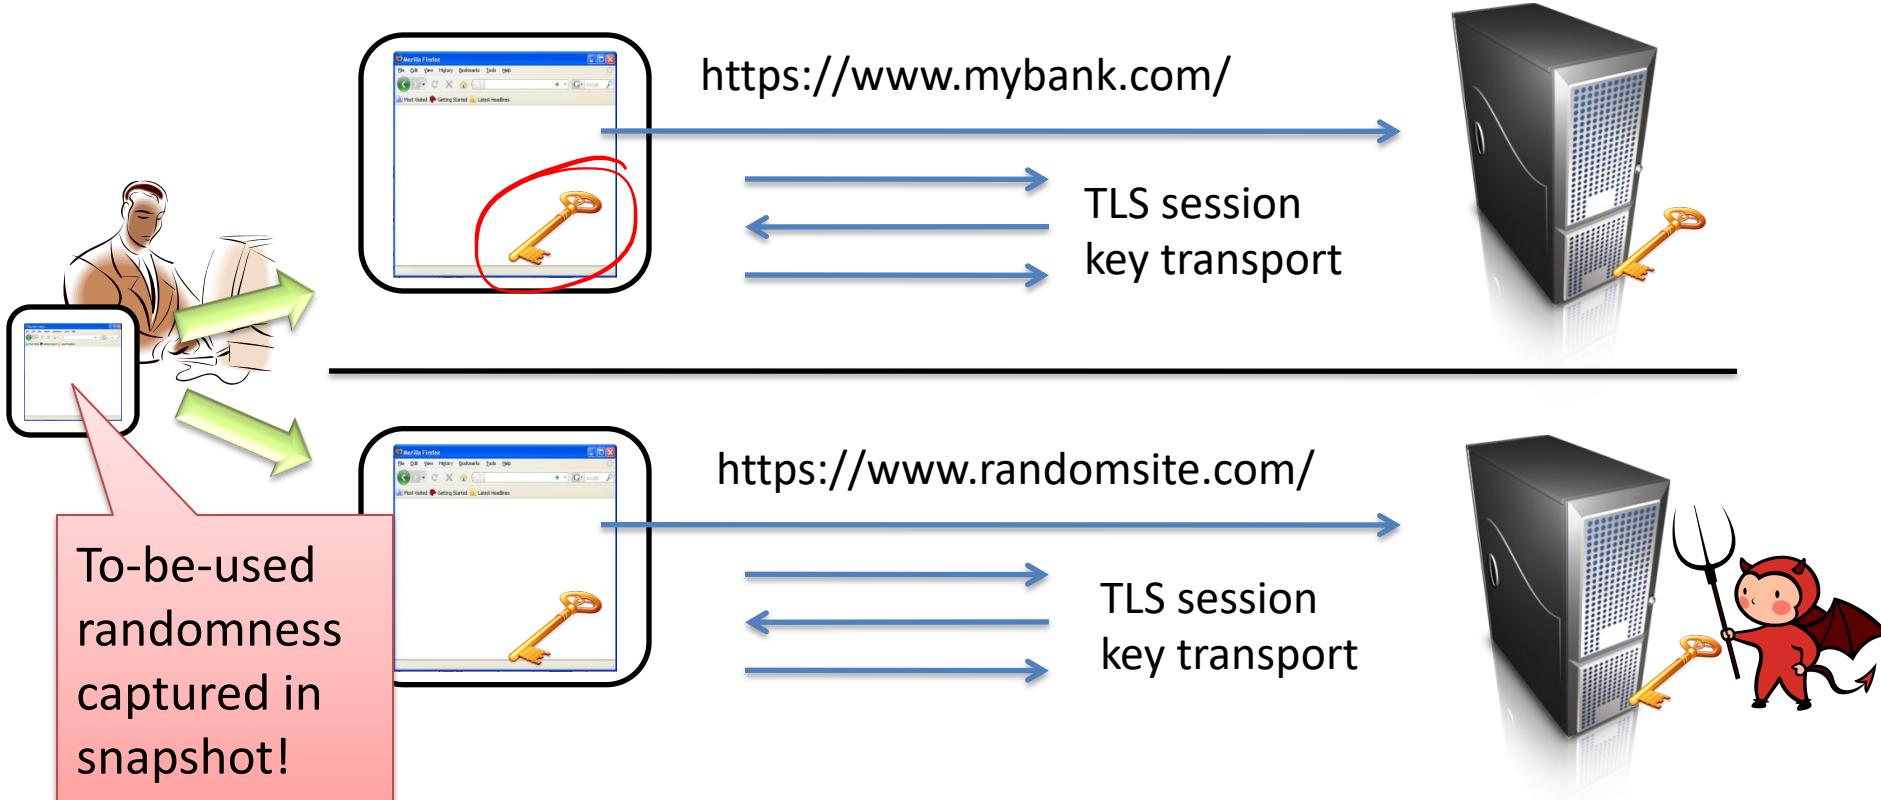
# Virtual machines and secure browsing

**“Protect Against Adware and Spyware:** Users protect their PCs against adware, spyware and other malware while browsing the Internet with Firefox in a virtual machine.”

[<http://www.vmware.com/company/news/releases/player.html>]



# Virtual machine resets lead to RNG failures for applications



Older versions of [Firefox](#), [Chrome](#) allow session compromise attacks

[Apache mod\\_ssl TLS server:](#)  
server's secret DSA key can be stolen!

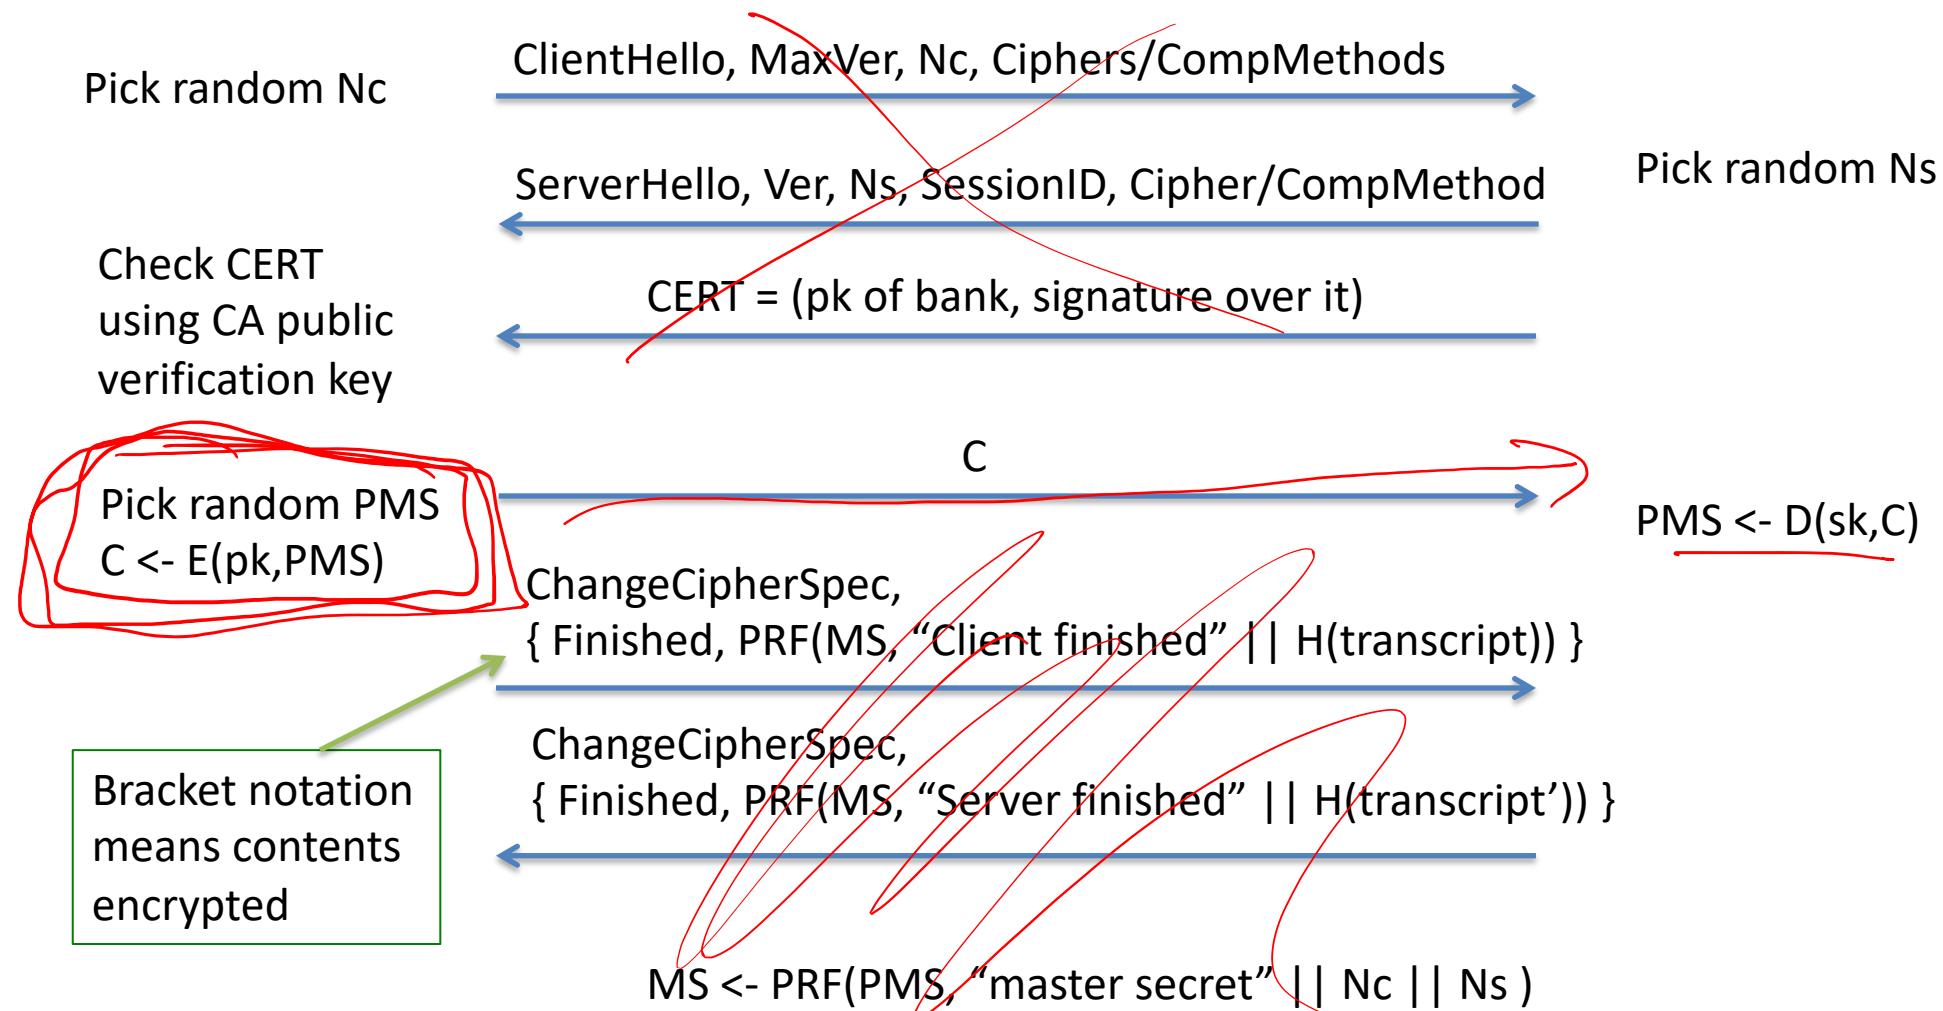


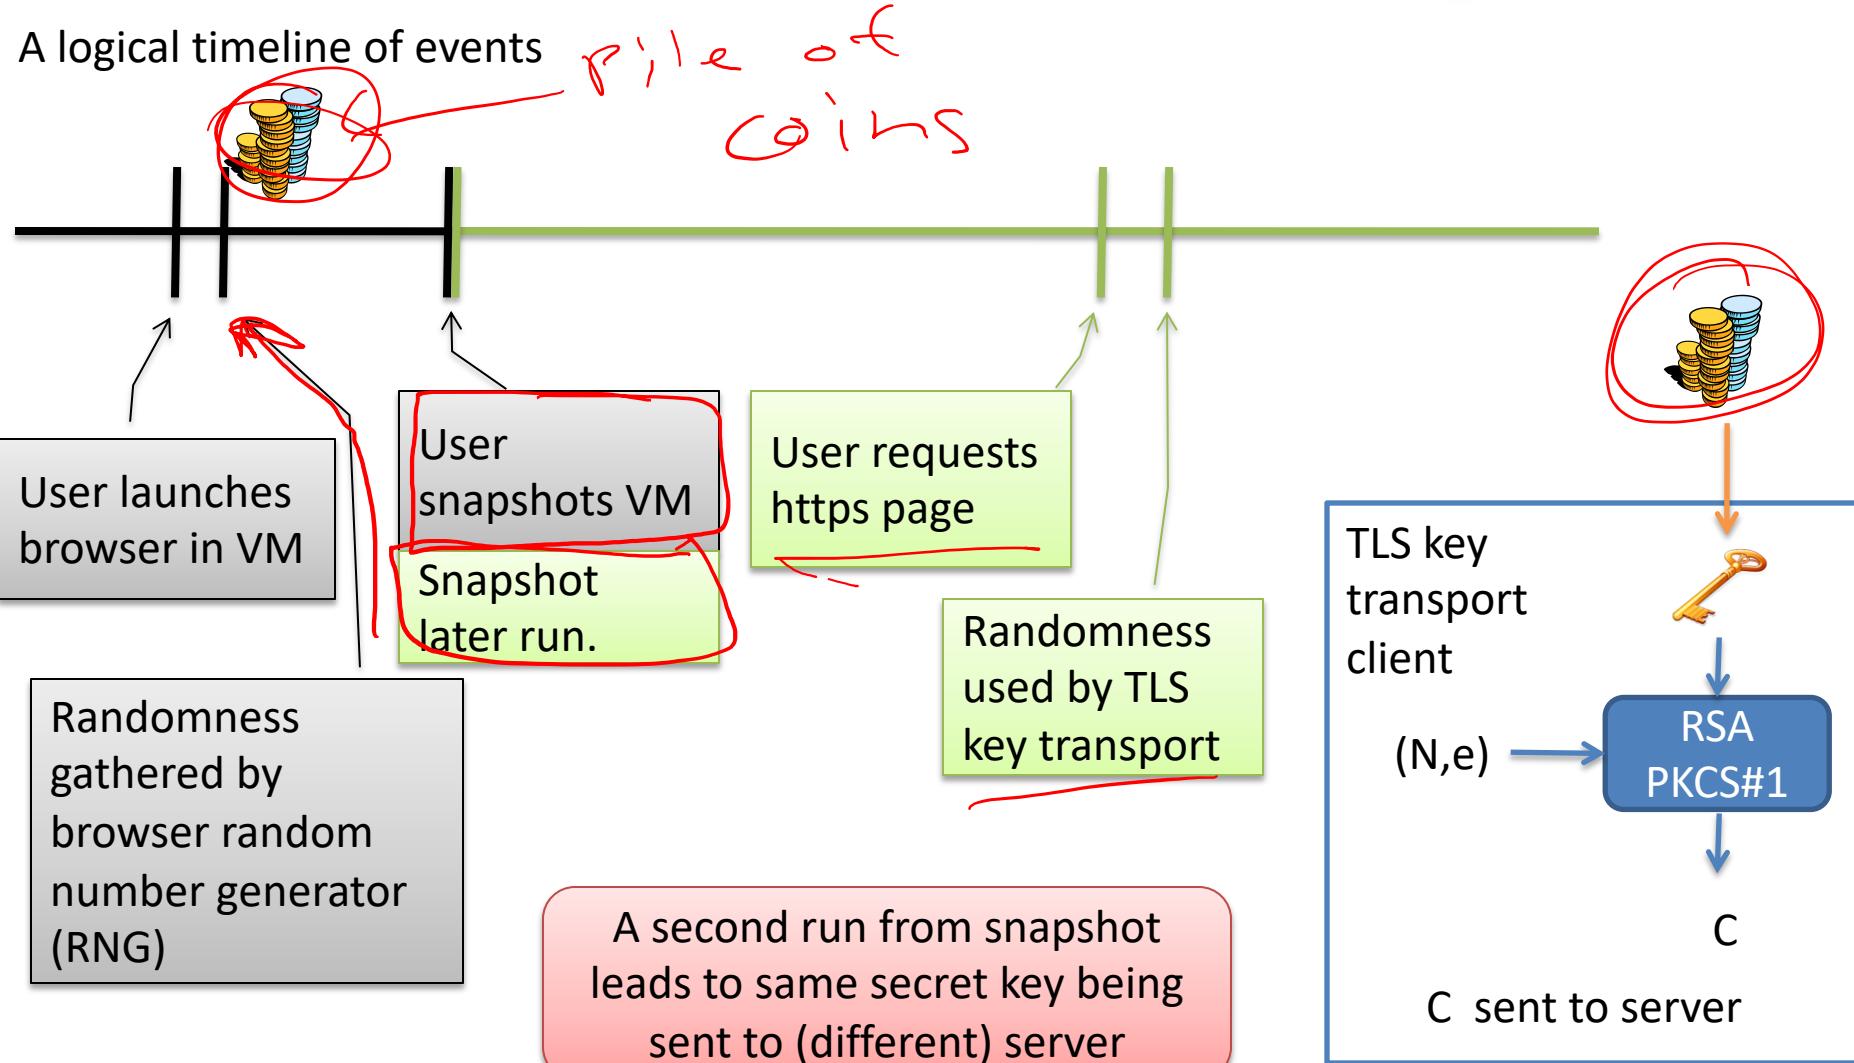
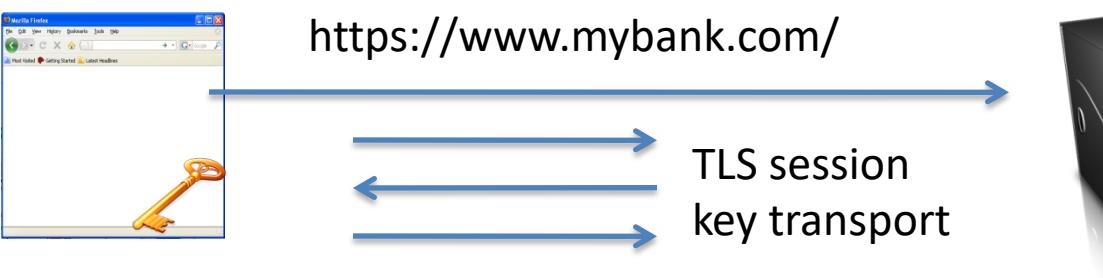
Client



Server

# TLS 1.2 handshake for RSA transport





# Reset vulnerabilities when using /dev/urandom after resumption?

- We showed that Linux /dev/(u)random and Windows system RNG are also vulnerable to resets
  - openssl genrsa will sometimes use repeat randomness (if ALSR is turned off, always)
- Primary problem is pooling structure of /dev/(u)random
- Changes have been made to Windows to fix

# Recent updates to Linux RNG

- [Donenfield 2022]
  - Removed complex pooling structure (called premature next attacks):
    - “As well, in the first place, our RNG never even properly handled the premature next issue, so removing an incomplete solution to a fake problem was particularly nice.” (<https://lore.kernel.org/lkml/20220522214457.37108-1-Jason@zx2c4.com/T/#u>)
  - Update cryptography (remove SHA1, replace with blake2s)
  - Different entropy accumulation, hashing strategy
  - Added better support for hypervisors to seed RNG on snapshot resumption
- Overall fantastic update and cleanup effort
- <https://www.zx2c4.com/projects/linux-rng-5.17-5.18/>

# Using RNGs

- Rule of thumb: more entropy is better
- In consuming applications:
  - Call cryptographically strong RNG such as /dev/urandom or Intel RDRAND
  - Can carefully mix in local sources of entropy if available, using hash tooling like HKDF
    - Hash it all together with cryptographic hash function to derive randomness to use
  - Minimize time between collection and use, when possible
- If efficiency is problem, use your own PRG seeded with above (be careful of reset vulnerabilities!)
  - Actually there's work on fixing this in Linux too, by having library support for fast user-space calls to /dev/random (



